

Evaluating eBPF as a Platform for Congestion Control Algorithm

Edward Wibowo Bokai Bi

Outline

- ▶ **Introduction**
- ▶ eBPFCCP Design
- ▶ Evaluation
- ▶ Results & Conclusions

Abstract

- ▶ **Goal:** Explore eBPF for Congestion Control Algorithms (CCAs).
- ▶ **Comparison:** Kernel-based (Linux TCP), eBPF-based, CCP, eBPFCCP.
- ▶ **Outcome:** eBPFCCP simplifies CCA deployment and development.

Introduction

- ▶ **Challenges with Kernel modules:**
 - ▶ Complex kernel programming.
 - ▶ Difficult adoption due to safety concerns.
 - ▶ Hard to get right.
- ▶ **eBPF Benefits:**
 - ▶ Verifier provides safety guarantees.
 - ▶ Has hooks for TCP congestion control.

Related Work

- ▶ **CCP:**
 - ▶ Congestion Control Plane.
 - ▶ Can implement CCAs independent of datapath.

eBPFCCP Design

▶ **Key Features:**

- ▶ User-space: `libccp` (from original CCP paper) for control logic and Rust layer for eBPF loading and management.
- ▶ Kernel-space: eBPF for flow hooks.
- ▶ Communication via BPF maps (BPF hash maps and BPF ring buffers).

eBPFCCP Design




Figure 1: Original CCP Design

eBPFCCP Design

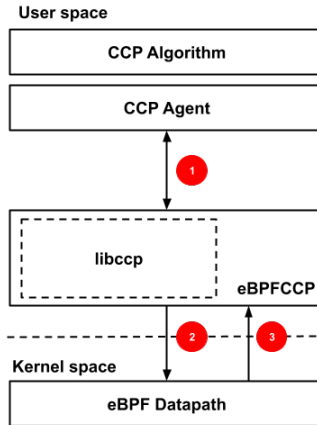


Figure 2: eBPFCCP Design

eBPFCCP Design

- ▶ **Goals:**
 - ▶ Match existing CCP behavior.
 - ▶ Minimize performance overhead.

Evaluation: Setup

- ▶ **Metrics:**
 - ▶ Behavior: Match existing CCAs.
 - ▶ Performance: CPU overhead.
- ▶ **Method:**
 - ▶ Use `iperf3` and `mahimahi`.
 - ▶ Controlled VM (single-core).

Results

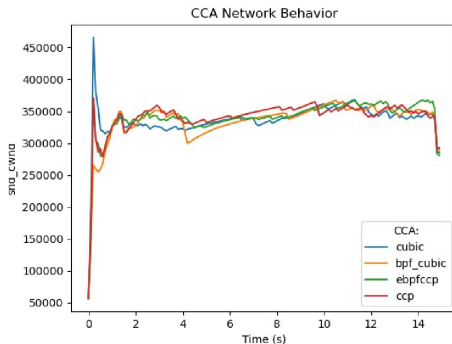


Figure 3: Behavior: Send Congestion Window Size Benchmark

Results




Figure 4: Behavior: RTT Benchmark

Results

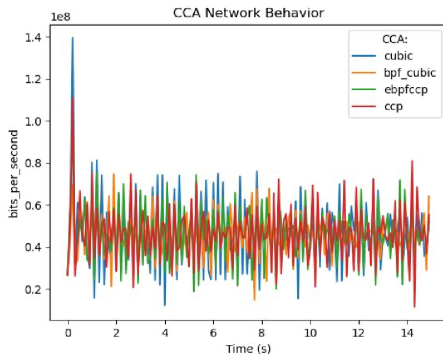


Figure 5: Behavior: Bits per Second Benchmark

Results




Figure 6: Performance Benchmark

Results

- ▶ **Behavior:**
 - ▶ All implementations behave similarly.
- ▶ **Performance:**
 - ▶ eBPF close to kernel-native.
 - ▶ eBPFCCP slightly slower due to communication overhead.

Limitations

- ▶ **Challenges:**

- ▶ Unix-pipe communication impacts performance.
- ▶ eBPF VM adds slight overhead.

- ▶ **Future Work:**

- ▶ Optimize communication channels.
- ▶ Integrate eBPFCCP directly with CCP.

Conclusion

- ▶ **eBPF as a Platform:**
 - ▶ Viable for safe, iterative CCA development.
 - ▶ Comparable performance to kernel-native CCAs.
- ▶ **Implications:**
 - ▶ Lowers entry barriers for research.
 - ▶ Promotes broader adoption.